

# Neural Networks as Parameterizers of Probability Distributions

ELG 5218 - Uncertainty Evaluation in Engineering Measurements and Machine Learning

Miodrag Bolic

University of Ottawa

February 25, 2026

- ① **Why NNs can parameterize distributions at all:** universal approximation  $\Rightarrow$  learn maps into parameter spaces.
- ② **General distribution parameterizations:**
  - mixtures / MDNs,
  - autoregressive factorizations,
  - latent-variable models (VAEs).
- ③ **Gaussian likelihood heads:** mean-variance networks (MVE), learning  $\sigma(x)$  without labels, failure modes and fixes.
- ④ **Where it is used:** a summary table.

## Step 1: Neural networks as function families

A feedforward neural network defines a parametric family of functions:

$$f_{\theta} : \mathbb{R}^d \rightarrow \mathbb{R}^m$$

with parameters  $\theta$  (weights/biases).

**Universal approximation:** with sufficient capacity, networks can approximate broad classes of functions to arbitrary accuracy (under standard conditions). (Cybenko 1989; Hornik 1991)

### Key implication

If a distribution can be described using (one or more) functions, then a NN can be used to approximate those functions, hence *parameterize* the distribution.

## Step 2: Constrained outputs $\Rightarrow$ valid distributions

To parameterize a distribution, NN outputs must respect constraints:

- probabilities sum to 1: use **softmax** (categorical / mixture weights),
- variances positive: use **exp** or **softplus**,
- covariance Symmetric Positive Definite SPD: e.g., Cholesky factorization.

So we define:

$$\eta(x) = g(f_\theta(x)),$$

where  $g(\cdot)$  enforces constraints and  $\eta(x)$  are valid distribution parameters.

# Two modeling goals: conditional vs. generative

**(A) Conditional density estimation:** model  $p_\theta(y | x)$

$$\eta(x) = f_\theta(x), \quad p_\theta(y | x) = p(y; \eta(x)).$$

**(B) Generative modeling:** model  $p_\theta(x)$ , often via latent noise  $z$ :

$$z \sim p(z), \quad x = \text{model}_\theta(z), \quad x \sim p_\theta(x).$$

Different families make  $p_\theta(\cdot)$  tractable in different ways.

# Generative modeling via latent variables

- Let  $Z$  be a *random noise variable*.

$$Z \sim p(z)$$

where  $p(z)$  is a simple prior (e.g. standard Gaussian or uniform).

- Let  $f_\theta$  be a *deterministic neural network*.

$$X_\theta = f_\theta(Z)$$

Since  $Z$  is random and  $f_\theta$  is deterministic,  $X_\theta$  is also a *random variable*.

- Therefore  $X_\theta$  has an induced distribution:

$$X_\theta \sim p_\theta(x).$$

## Key idea:

- $p_\theta(x)$  is *not* output by the neural network.
- Instead,  $p_\theta(x)$  is the *distribution of samples* produced when noise  $Z \sim p(z)$  is transformed through the network  $f_\theta$ .
- In other words: the NN *defines* the distribution by pushing  $p(z)$  through  $f_\theta$ .

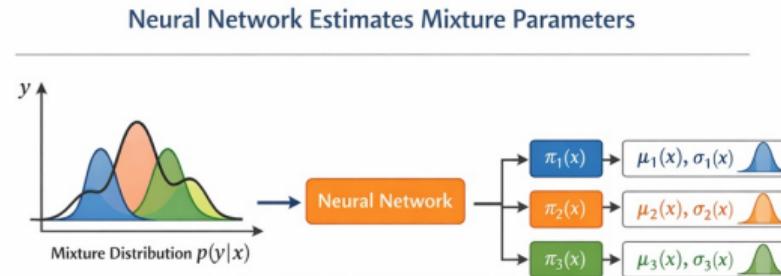
# Family 1: Mixtures / Mixture Density Networks (MDNs)

**Idea:** Approximate complex (multi-modal) densities using mixtures of simple components.  
Conditional mixture model:

$$p_{\theta}(y | x) = \sum_{k=1}^K \pi_k(x) p_k(y; \theta_k(x)),$$

where a NN outputs  $\{\pi_k(x), \theta_k(x)\}_{k=1}^K$ .

MDNs were introduced as NN-controlled mixture models to represent *arbitrary conditional densities* in principle. (Bishop 1994)



## Family 2: Autoregressive factorization (e.g., MADE)

**Key theorem (probability chain rule):**

$$p(x) = \prod_{d=1}^D p(x_d | x_{<d}).$$

**Neural autoregressive density estimation:** a NN outputs parameters of each conditional factor while enforcing the dependency structure.

MADE uses masking to guarantee that output  $d$  depends only on  $x_{<d}$  and yields tractable likelihood. (Germain et al. 2015)

# Autoregressive Model of Order 2 (AR(2))

A classical AR(2) process specifies that the next value depends on the two most recent past values:

$$x_t = a_1 x_{t-1} + a_2 x_{t-2} + \varepsilon_t, \quad \varepsilon_t \sim \mathcal{N}(0, \sigma^2).$$

This implies a conditional density:

$$p(x_t | x_{t-1}, x_{t-2}) = \mathcal{N}(a_1 x_{t-1} + a_2 x_{t-2}, \sigma^2).$$

**Key idea:** autoregression factorizes the sequence density using the probability chain rule:

$$p(x_{1:T}) = p(x_1) p(x_2 | x_1) \prod_{t=3}^T p(x_t | x_{t-1}, x_{t-2}).$$

# Neural Generalization of AR(2)

Replace the linear mapping with a neural network:

$$x_t \mid x_{t-1}, x_{t-2} \sim \mathcal{N}\left(\mu_\theta(x_{t-1}, x_{t-2}), \sigma_\theta(x_{t-1}, x_{t-2})^2\right).$$

## Neural networks:

- Input:  $(x_{t-2}, x_{t-1})$
- Output: mean  $\mu_\theta$  and log-std  $\log \sigma_\theta$

Example tiny NNs:

$$\mu_\theta : \mathbb{R}^2 \rightarrow \mathbb{R}, \quad 2 \rightarrow 8 \rightarrow 1,$$

$$\log \sigma_\theta : \mathbb{R}^2 \rightarrow \mathbb{R}, \quad 2 \rightarrow 4 \rightarrow 1.$$

**Why useful:** NN can model nonlinear dynamics, heteroscedastic noise, and complex temporal patterns that linear AR models cannot capture.

# Training Data for AR(2) Neural Networks

Given a time series  $x_1, x_2, \dots, x_T$ :

**Construct training triples:**

$$(x_{t-2}, x_{t-1}) \rightarrow x_t, \quad t = 3, \dots, T.$$

This yields a dataset:

$$\{((x_{t-2}, x_{t-1}), x_t)\}_{t=3}^T.$$

**Training objective:** Gaussian negative log-likelihood

$$\ell_t = \log \sigma_\theta(x_{t-1}, x_{t-2}) + \frac{1}{2} \left( \frac{x_t - \mu_\theta(x_{t-1}, x_{t-2})}{\sigma_\theta(x_{t-1}, x_{t-2})} \right)^2.$$

Total loss:

$$\mathcal{L}(\theta) = \sum_{t=3}^T \ell_t.$$

**Optimization:** gradient descent on  $\mathcal{L}(\theta)$ .

# Intuition: What the AR(2) Neural Network Learns

**Inputs to the networks:**

$$x_{t-2}, x_{t-1}$$

**Outputs:**

$$\mu_t, \sigma_t$$

The NN learns:

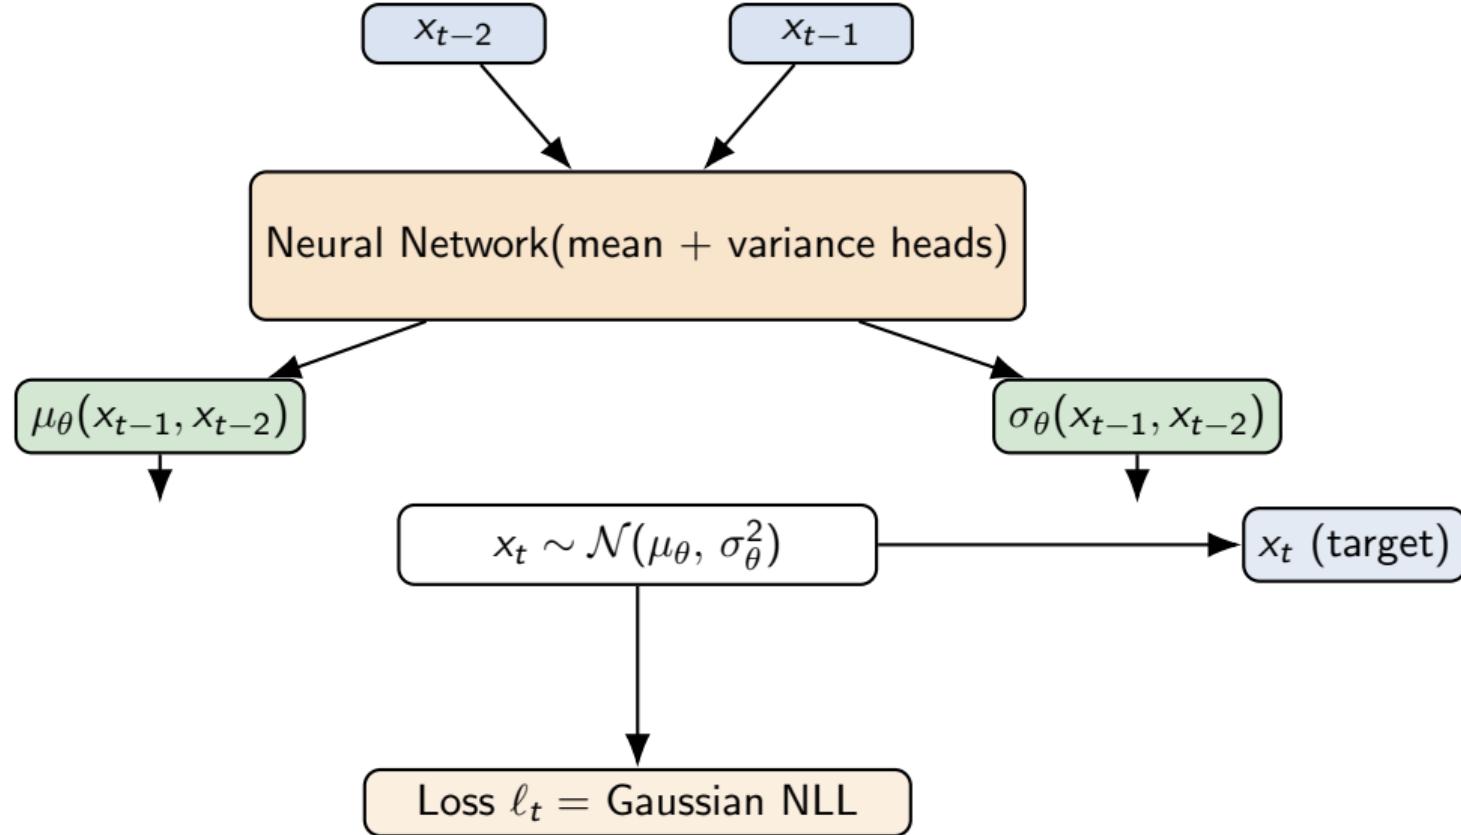
- nonlinear temporal relationships,
- how uncertainty  $\sigma_t$  varies with past states (heteroscedasticity),
- how to represent complex dynamics beyond linear AR(2).

**Sampling:**

$$x_t \sim \mathcal{N}(\mu_\theta(x_{t-1}, x_{t-2}), \sigma_\theta^2(x_{t-1}, x_{t-2})).$$

**Full sequence generation:** draw  $x_1, x_2$  from base distributions, then sequentially generate  $x_3, \dots, x_T$  using NN conditionals.

# AR(2) Neural Autoregressive Model



## Family 3: Latent-variable models (VAEs)

**Generative story:**

$$z \sim p(z), \quad x | z \sim p_\theta(x | z)$$

where  $p_\theta(x | z)$  is parameterized by a NN (decoder).

**Inference is hard:** posterior  $p_\theta(z | x)$  is usually intractable, so VAEs introduce

$$q_\phi(z | x)$$

(encoder NN) and optimize the ELBO. (Kingma and Welling 2013)

## Gaussian head = simplest conditional distribution parameterization

Now specialize to a simple conditional family:

$$y | x \sim \mathcal{N}(\mu_\theta(x), \sigma_\theta(x)^2).$$

A neural network outputs:

$$(\mu_\theta(x), \log \sigma_\theta(x)).$$

This is the classic Mean–Variance Estimation (MVE) network formulation, dating back to early work on learning mean and variance from data via likelihood. (Nix and Weigend 1994)

# Gaussian NLL and why $\sigma(x)$ is learnable without labels

Gaussian negative log-likelihood (per datum), up to constants:

$$\mathcal{L}(x, y) = \log \sigma_\theta(x) + \frac{1}{2} \left( \frac{y - \mu_\theta(x)}{\sigma_\theta(x)} \right)^2.$$

Why this provides “supervision” for variance

The loss penalizes:

- large uncertainty ( $\log \sigma$  term),
- underestimating uncertainty when residuals are large (scaled residual term).

This likelihood-based learning of variance without variance targets is explicitly derived in classic MVE work. (Nix and Weigend 1994)

## Interpretation: optimal $\sigma$ tracks conditional residual scale

Fix  $\mu(x)$  and optimize  $\sigma$  for residual  $r = y - \mu(x)$ :

$$\ell(\sigma) = \log \sigma + \frac{r^2}{2\sigma^2}.$$

Stationary point satisfies:

$$\frac{d\ell}{d\sigma} = 0 \Rightarrow \sigma^2 = r^2.$$

Over many samples (conceptually at similar  $x$ ):

$$\sigma(x)^2 \approx \mathbb{E}[(y - \mu(x))^2 | x].$$

This explains why variance can be learned from data scatter even without explicit variance labels. (Nix and Weigend 1994)

## Pathology: $\sigma(x)$ can “explain away” mean error

Because the NLL trades off mean fit and variance:

$$\log \sigma(x) + \frac{(y - \mu(x))^2}{2\sigma(x)^2},$$

in regions where  $\mu_\theta(x)$  is imperfect, increasing  $\sigma_\theta(x)$  can reduce loss.

Recent work highlights that likelihood gradients can be *scaled* by predictive variance, potentially compromising mean learning. (Immer et al. 2023)

Predicted  $\sigma(x)$  becomes structured (e.g., follows the signal) even if true noise is constant.

# Training practice: mean warm-up and regularization

Classic and modern guidance for MVE-style networks includes:

- **Warm-up:** fit mean first with fixed variance; then train mean+variance jointly.
- **Regularize** variance or its smoothness to prevent spurious structure.

These issues and recommendations are discussed explicitly in recent work on optimal training of MVE networks. (Sluijterman et al. 2023)

# Training Algorithm for Gaussian-Head (MVE) Networks

## Algorithm: Stable Training with Mean Warm-Up

### Inputs:

- Dataset  $\mathcal{D} = \{(x_i, y_i)\}_{i=1}^N$
- Network  $f_\theta(x)$  producing  $(\mu_\theta(x), \log \sigma_\theta(x))$
- Warm-up epochs  $T_{\text{warm}}$
- Regularization weight  $\lambda$

## Phase 1: Mean Warm-Up

- ① Freeze variance parameters (or set  $\sigma_\theta(x) = \text{const}$ ).

- ② For  $t = 1, \dots, T_{\text{warm}}$ :

- For each minibatch  $B$ :
- Compute  $\mu_\theta(x)$  for all  $(x, y) \in B$ .
- Use fixed-variance loss:

$$\ell_i = \frac{(y_i - \mu_\theta(x_i))^2}{2\sigma_{\text{fixed}}^2}.$$

- Update only  $\theta_\mu$ .

## Phase 2: Joint Mean + Variance Training

- ③ Unfreeze variance parameters.
- ④ For  $t = T_{\text{warm}} + 1, \dots, T_{\text{total}}$ :
  - Compute  $(\mu_\theta(x), \log \sigma_\theta(x))$  for all  $(x, y) \in B$ .
  - Compute  $\sigma_\theta(x) = \exp(\log \sigma_\theta(x))$ .
  - Compute NLL:

$$\ell_i = \log \sigma_\theta(x_i) + \frac{(y_i - \mu_\theta(x_i))^2}{2\sigma_\theta(x_i)^2}.$$

- Add variance regularization:

$$R = \lambda \|\log \sigma_\theta(x)\|^2 \quad \text{or} \quad \lambda \|\nabla_x \sigma_\theta(x)\|^2.$$

- Update  $(\theta_\mu, \theta_\sigma)$  using loss  
 $L = \frac{1}{|B|} \sum_i \ell_i + R$ .

# How do we evaluate without “uncertainty labels”?

True uncertainty targets rarely exist. Evaluation uses:

- **Held-out NLL / log predictive density** (proper scoring rule),
- **Calibration** checks (coverage, standardized residuals),
- Robustness under dataset shift (OOD uncertainty).

Deep ensembles are a widely-used baseline emphasizing calibration and log scores for uncertainty evaluation. (Lakshminarayanan et al. 2017)

# Calibration test for Gaussian heads

Standardized residuals:

$$z_i = \frac{y_i - \mu_\theta(x_i)}{\sigma_\theta(x_i)}.$$

If the model is well-specified and calibrated:

$$z \sim \mathcal{N}(0, 1).$$

Coverage checks:

- $\Pr(|y - \mu| \leq 1\sigma) \approx 0.68$
- $\Pr(|y - \mu| \leq 2\sigma) \approx 0.95$

These are the practical diagnostics that complement NLL when uncertainty labels are absent.  
(Lakshminarayanan et al. 2017)

# Where NN distribution parameterization is used (summary)

Family	Distribution form	NN outputs	Tractability	Typical uses
Gaussian head (MVE) (Nix and Weigend 1994)	$\mathcal{N}(\mu(x), \sigma(x)^2)$	$\mu(x), \log \sigma(x)$	Exact NLL	Probabilistic regression, aleatoric uncertainty
MDN / mixtures (Bishop 1994)	$\sum_k \pi_k(x) p_k(y)$	$\pi_k, \theta_k$	Exact NLL	Multi-modal $p(y x)$ , inverse problems
Autoregressive (MADE) (Germain et al. 2015)	$\prod_d p(x_d   x_{<d})$	params of each conditional	Exact LL	High-dim density estimation, compression
VAE (Kingma and Welling 2013)	$p(x)$	encoder+decoder params	ELBO (approx.)	Representation learning, generation

# Take-home messages

- NNs can parameterize distributions because they approximate functions; with constraint mappings, outputs become valid distribution parameters. (Cybenko 1989; Hornik 1991)
- “Universal” density constructions come from mixtures, autoregressive factorizations, flows, latent-variable models, and EBMs. (Bishop 1994; Germain et al. 2015; Papamakarios et al. 2021; Kingma and Welling 2013; LeCun et al. 2006)
- Gaussian mean–variance networks are a simple special case;  $\sigma(x)$  is learned indirectly by likelihood (no variance labels needed). (Nix and Weigend 1994)
- Joint mean/variance training can produce pathologies (variance absorbing mean error); warm-up, regularization, and careful calibration matter. (Immer et al. 2023; Stirn et al. 2023; Sluijterman et al. 2023)

# References I

-  Bishop, Christopher M. (1994). *Mixture Density Networks*. Neural Computing Research Group Technical Report NCRG/94/004. Aston University. URL:  
<https://eprints.aston.ac.uk/373/>.
-  Cybenko, George (1989). "Approximation by Superpositions of a Sigmoidal Function". In: *Mathematics of Control, Signals and Systems* 2.4, pp. 303–314. URL:  
<http://aaclab.mit.edu/material/handouts/universalapproximation-1989.pdf>.
-  Germain, Mathieu et al. (2015). "MADE: Masked Autoencoder for Distribution Estimation". In: *Proceedings of Machine Learning Research* 37, pp. 881–889. URL:  
<https://proceedings.mlr.press/v37/germain15.html>.
-  Hornik, Kurt (1991). "Approximation Capabilities of Multilayer Feedforward Networks". In: *Neural Networks* 4.2, pp. 251–257. DOI: 10.1016/0893-6080(91)90009-T. URL:  
[https://web.njit.edu/~usman/courses/cs677\\_spring21/hornik-nn-1991.pdf](https://web.njit.edu/~usman/courses/cs677_spring21/hornik-nn-1991.pdf).

## References II

- Immer, Alexander et al. (2023). "Effective Bayesian Heteroscedastic Regression with Deep Neural Networks". In: *Advances in Neural Information Processing Systems*. URL: [https://papers.nips.cc/paper\\_files/paper/2023/file/a901d5540789a086ee0881a82211b63d-Paper-Conference.pdf](https://papers.nips.cc/paper_files/paper/2023/file/a901d5540789a086ee0881a82211b63d-Paper-Conference.pdf).
- Kingma, Diederik P. and Max Welling (2013). "Auto-Encoding Variational Bayes". In: *arXiv preprint arXiv:1312.6114*. URL: <https://arxiv.org/abs/1312.6114>.
- Lakshminarayanan, Balaji, Alexander Pritzel, and Charles Blundell (2017). "Simple and Scalable Predictive Uncertainty Estimation using Deep Ensembles". In: *Advances in Neural Information Processing Systems*. URL: <https://arxiv.org/abs/1612.01474>.
- LeCun, Yann et al. (2006). "A Tutorial on Energy-Based Learning". In: Tutorial chapter in *Predicting Structured Data* (MIT Press). URL: <http://yann.lecun.com/exdb/publis/pdf/lecun-06.pdf>.

## References III

-  Nix, David A. and Andreas S. Weigend (1994). "Estimating the Mean and Variance of the Target Probability Distribution". In: *Proceedings of the 1994 IEEE International Conference on Neural Networks*, pp. 55–60. DOI: 10.1109/ICNN.1994.374138. URL: <https://ieeexplore.ieee.org/document/374138>.
-  Papamakarios, George et al. (2021). "Normalizing Flows for Probabilistic Modeling and Inference". In: *Journal of Machine Learning Research* 22.57, pp. 1–64. URL: <https://jmlr.org/papers/volume22/19-1028/19-1028.pdf>.
-  Sluijterman, Laurens, Eric Cator, and Tom Heskes (2023). "Optimal Training of Mean Variance Estimation Neural Networks". In: *arXiv preprint arXiv:2302.08875*. URL: <https://arxiv.org/pdf/2302.08875v1>.
-  Stirn, Andrew et al. (2023). "Faithful Heteroscedastic Regression with Neural Networks". In: *Proceedings of Machine Learning Research* 206. URL: <https://proceedings.mlr.press/v206/stirn23a/stirn23a.pdf>.